Runtime Execution Governance for AI Systems: A Cross-Platform Synthesis and Architectural Framework

John M. Willis
Sustainable Future Tech
jwillis@sustainablefuturetech.com
ORCID: 0009-0007-5889-9628

March 30, 2026

Abstract

AI governance is undergoing a structural transition from policy-based, pre-deployment oversight to execution-time control of autonomous and distributed systems. While prior work has addressed policy frameworks, guardrails, and monitoring, a coherent architecture that unifies governance intent, execution semantics, and system realization remains absent.

This paper presents a cross-platform synthesis of research, engineering practice, enterprise systems, regulatory frameworks, and practitioner discourse, identifying a converging paradigm: governance as a runtime decision function over system state transitions. Building on this synthesis, we introduce a vertically integrated governance stack consisting of governance compilation, the Runtime Governance Architecture (RGA), the Artificial Intelligence Governance Control Plane (AGCP), and the PBSAI governance ecosystem.

RGA formalizes governance as a deterministic execution model over state transitions; AGCP provides a control-plane implementation that enforces execution-bound authorization, lifecycle derivation, and invariant correctness; and PBSAI operationalizes governance as a domain-structured, evidence-centric multi-agent system. Governance compilation bridges policy intent to executable artifacts across these layers.

We further provide a comparative systems analysis demonstrating that existing approaches—including policy engines, guardrails frameworks, and regulatory models—lack end-to-end alignment between intent, enforcement, and execution. The proposed stack resolves this gap by enabling deterministic, auditable, and non-bypassable governance at runtime.

The results position runtime governance not as an extension of policy enforcement, but as a foundational architectural layer for AI-enabled systems.

Keywords: AI governance, runtime governance, control plane, multi-agent systems, policy enforcement, deterministic systems, governance compilation, explainable AI, auditability, distributed systems

Cite as: Willis, J. M. (2026). *Runtime Execution Governance for AI Systems: A Cross-Platform Synthesis and Architectural Framework.* Sustainable Future Tech, Inc.
https://doi.org/10.5281/zenodo.19341177

1 Introduction

The emergence of agentic AI systems and autonomous execution pipelines has exposed fundamental limitations in traditional governance models. Existing approaches emphasize policy definition, compliance processes, and post hoc auditing. However, these mechanisms do not provide control over system behavior at the point where it matters most: the moment of execution.

Recent developments across research, engineering, and industry indicate a shift toward runtime governance, where decisions must be evaluated and enforced in real time as systems act. Despite this convergence, current approaches remain fragmented. Policy frameworks are typically non-executable, enforcement mechanisms are platform-specific and inconsistent, and runtime systems often operate without a unified governance model.

This paper advances the thesis that:

> AI governance is fundamentally a runtime control problem over state transitions.

To address this, we synthesize perspectives across multiple ecosystems and introduce a vertically integrated governance stack that connects policy intent to execution-time enforcement.

This stack consists of:

- **Runtime Governance Architecture (RGA)**, which defines governance as a deterministic state-transition execution model in which governance intent is compiled into canonical artifacts and evaluated through a consistent governance execution pipeline (Willis, 2026d);
- the **Artificial Intelligence Governance Control Plane (AGCP)**, which enforces deterministic, execution-bound authorization and lifecycle correctness (Willis, 2026a); and
- **The Practitioner's Blueprint for Secure AI (PBSAI)**, a domain-structured, evidence-centric multi-agent architecture that operationalizes governance within AI estates (Willis, 2026e).

Together, these components form a coherent architecture that enables governance to be expressed, executed, and verified at runtime.

Figure 1 summarizes this integrated governance stack. It shows how governance intent is transformed into executable artifacts, interpreted through a runtime governance model, enforced by a deterministic control plane, and ultimately realized within a domain-structured multi-agent AI estate.

This paper builds upon prior work introducing the Runtime Governance Architecture (RGA), the Artificial Intelligence Governance Control Plane (AGCP), and the PBSAI governance ecosystem (Willis, 2026d;a;e).

The remainder of this paper proceeds as follows. We first synthesize current thinking across platforms, then introduce the integrated governance stack, provide a comparative systems analysis, and conclude with implications for future AI system design.

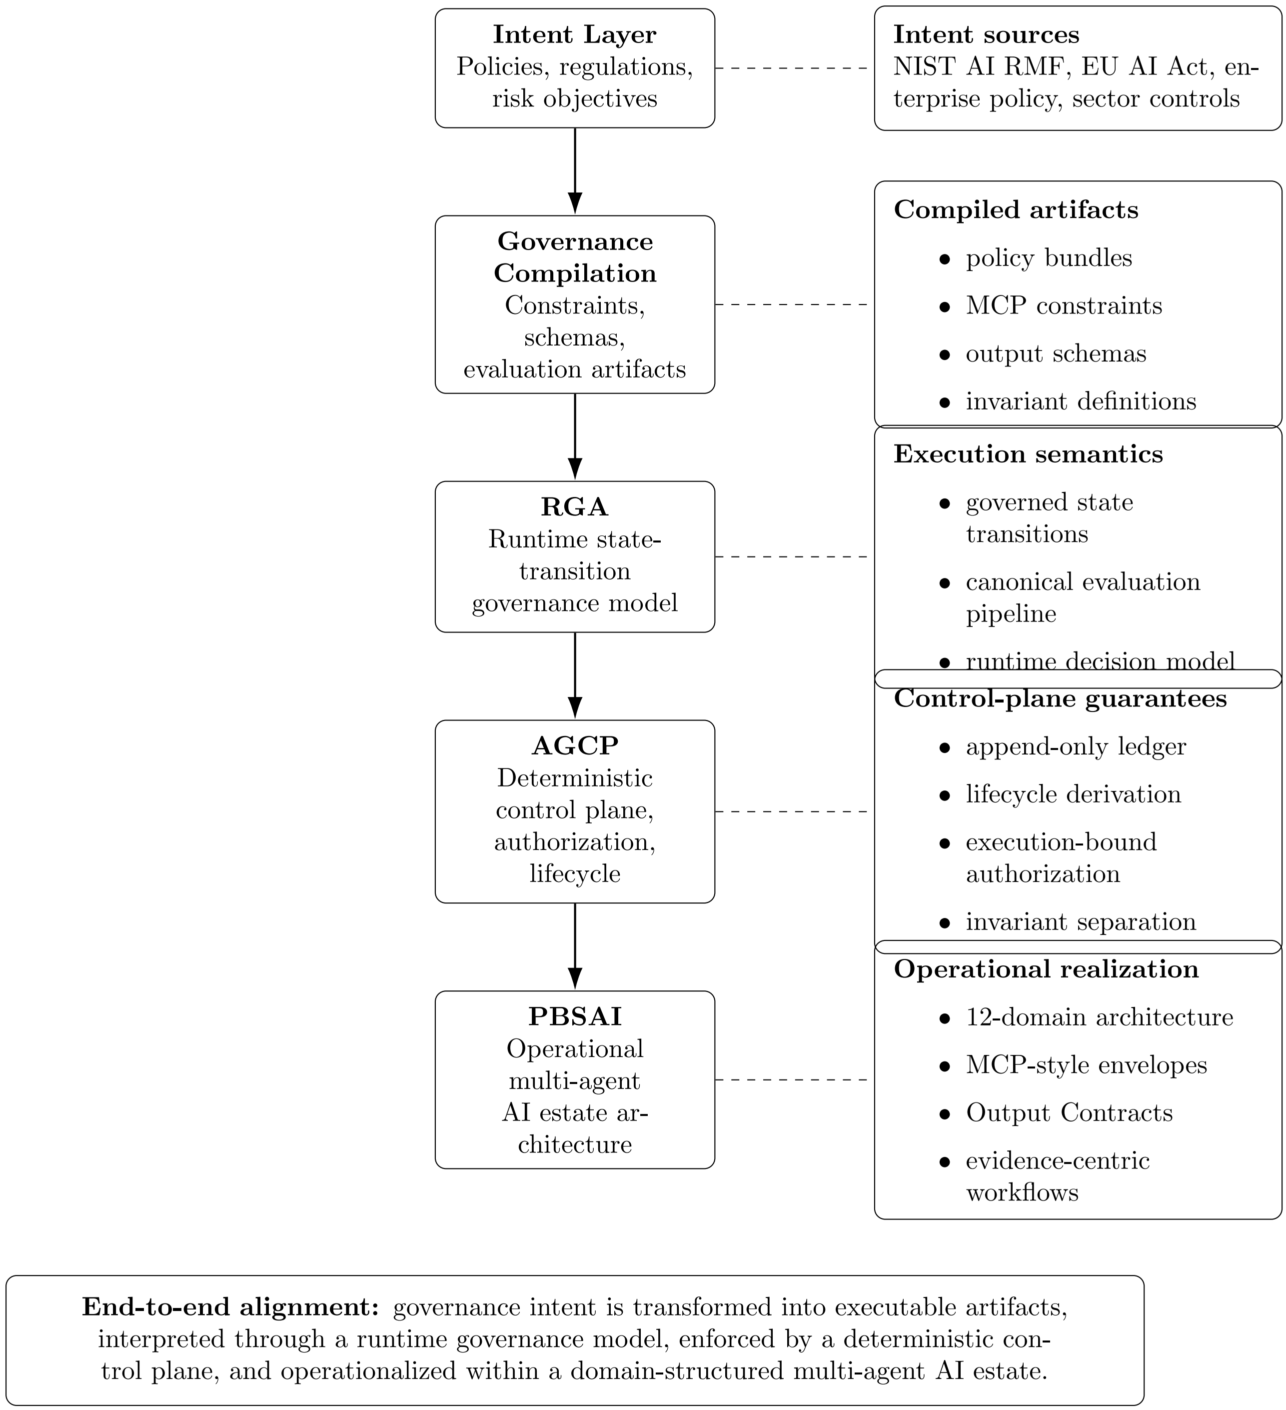

Figure 1: Integrated governance stack showing the alignment of governance intent, governance compilation, the Runtime Governance Architecture (RGA), the Artificial Intelligence Governance Control Plane (AGCP), and the PBSAI operational architecture. PBSAI provides the realized multi-agent system structure; RGA defines the runtime governance model; AGCP enforces deterministic execution-time correctness; and governance compilation bridges human-readable policy intent to executable governance artifacts.

2 Contributions

This paper makes the following contributions:

1. **Cross-platform synthesis:** We integrate perspectives from research literature, engineering systems, enterprise practices, regulatory frameworks, and practitioner discourse to characterize the current state of AI governance.
2. **Runtime governance framing:** We formalize the emerging paradigm of governance as a runtime decision function over state transitions, clarifying its implications for system design.
3. **Governance compilation abstraction:** We introduce governance compilation as the transformation layer that converts policy intent into executable artifacts, bridging the gap between human-readable governance and machine-executable control.
4. **Integrated governance stack:** We present a unified architecture linking governance compilation, RGA (Willis, 2026d), AGCP (Willis, 2026a), and PBSAI (Willis, 2026e), enabling end-to-end alignment from policy intent to runtime enforcement.
5. **Comparative systems analysis:** We demonstrate that existing approaches, including policy engines, guardrails frameworks, and regulatory models, lack full-stack alignment, and we show how the proposed architecture resolves these gaps.
6. **Architectural positioning:** We establish runtime governance as a foundational control-plane concern, rather than an extension of policy or monitoring systems.

3 Research Perspective: Governance as a Function

Recent research reframes governance as a decision function applied during execution. Work such as Gaurav et al. (2025) proposes governance layers that intercept and evaluate actions in real time. Similarly, Mazzocchetti (2026) formalizes policy constraints as execution conditions enforced through runtime validation and audit mechanisms.

Agent safety research further supports this shift. Empirical studies of agent behavior demonstrate that static evaluation is insufficient; safety must be assessed in tool-using, multi-step environments where actions unfold over time.

These developments imply a functional model:

$$f(s, a, c) \to d \tag{1}$$

where:

- s = system state
- a = proposed action
- c = context
- d = decision (allow, constrain, escalate, deny)

4 Implementation Perspective: Governance as a Control Plane

In engineering practice, runtime governance is emerging as a control plane separating agent logic from action execution.

The Microsoft Agent Governance Toolkit defines governance infrastructure as operating between agent frameworks and executed actions, enabling deterministic enforcement, identity control, and sandboxing (Microsoft, 2025). This enables deterministic policy enforcement, zero-trust identity, and sandboxing.

Similarly, runtime security architectures emphasize real-time defense against adversarial manipulation (Microsoft, 2026).

These systems implement governance through:

- action interception
- policy evaluation
- execution gating
- audit logging

5 Enterprise Systems: From Monitoring to Enforcement

Enterprise AI systems are transitioning from monitoring toward enforcement.

OpenAI emphasizes guardrails that constrain behavior during operation (OpenAI, 2024). Anthropic embeds normative constraints directly within models through Constitutional AI (Bai et al., 2022).

Google DeepMind emphasizes scalable oversight and defense-in-depth (DeepMind, 2024), while Meta focuses on layered safeguards and responsible deployment (Meta, 2024).

However, these approaches remain partially aligned with runtime governance and lack a unified execution-layer architecture.

6 Regulatory Perspective: Monitoring Without Mediation

Regulatory frameworks emphasize lifecycle governance rather than execution control.

The NIST AI Risk Management Framework defines governance across stages of system development and deployment (NIST, 2023). The EU AI Act mandates logging and post-market monitoring but does not define real-time enforcement mechanisms (EU, 2024).

Similarly, ENISA and CISA frame governance in terms of cybersecurity and operational risk (ENISA, 2023; CISA, 2024).

These frameworks support observability but not action mediation.

7 Public and Practitioner Discourse

Public discourse reveals uneven adoption.

YouTube engineering content increasingly frames governance as execution-time control (YouTube, 2025). In contrast, practitioner discussions on Reddit focus on compliance and tooling rather than runtime enforcement (Reddit, 2025).

This divergence highlights a gap between conceptual advances and practical understanding.

8 Industry and Market Trends

Industry analysis recognizes governance as an emerging product category.

Gartner identifies AI governance platforms as enabling automated policy enforcement and anomaly detection (Gartner, 2026). McKinsey notes that organizations are only beginning to adapt governance for agentic systems (McKinsey, 2025). Enterprise and consulting guidance typically emphasizes governance structures, risk management processes, and organizational operating models. These approaches remain largely interpretive and are not expressed as executable, runtime-enforceable artifacts. Governance that is not expressed as a deterministic, executable artifact cannot provide guarantees about system behavior at runtime.

9 Incident-Driven Evidence

Media coverage of AI failures consistently points to runtime failures.

Reports of rogue agents and AI security risks highlight failures occurring during execution rather than design (TechCrunch, 2026).

10 Alignment and Safety Communities

Alignment-focused communities provide empirical evidence that agents do not reliably adhere to safety constraints under real-world conditions.

Testing shows failures of rule adherence under task pressure (LessWrong, 2025), reinforcing the need for execution-time governance.

11 Alignment with Runtime Governance Architecture and AGCP

As shown in Figure 1, the PBSAI Governance Ecosystem, the Runtime Governance Architecture (RGA), and the Artificial Intelligence Governance Control Plane (AGCP) form a vertically integrated governance stack.

PBSAI provides the operational realization layer, structuring the AI estate as a domain-organized, multi-agent system with standardized context envelopes and evidence contracts (Willis, 2026e). RGA defines governance as a deterministic state-transition execution model in which governance intent is compiled into canonical artifacts and evaluated through a consistent execution pipeline (Willis, 2026d). AGCP enforces this model at runtime through deterministic, execution-bound authorization and lifecycle derivation (Willis, 2026a), with normative interfaces defined in the AGCP Specification (Willis, 2026b).

The integrated governance stack can be summarized as:

- **Intent layer**: Policies, regulations, and organizational constraints define desired system properties.
- **Execution model (RGA)**: Governance intent is compiled into executable artifacts and evaluated as governed state transitions.
- **Control plane (AGCP)**: Deterministic validation, authorization, and lifecycle derivation are enforced at execution time.
- **Operational layer (PBSAI)**: The AI estate is instantiated as a multi-agent system producing structured, auditable evidence.

Within this stack, PBSAI does not enforce governance correctness; rather, it ensures that all system activity is expressed in a form that can be evaluated and governed deterministically. RGA defines the execution model, and AGCP enforces its correctness.

11.1 Formal Execution Semantics (Informal)

The governance execution model is grounded in a formal execution semantics (Willis, 2026c). Systems are modeled as state-transition processes in which proposed actions are validated prior to execution.

Validation evaluates each proposed action against:

- **Compiled constraints**, derived from governance policies; and
- **System invariants**, which must hold over all system states.

A transition is authorized only if both conditions are satisfied:

$$\mathrm{VALID}(S, A) = \mathrm{true}$$

where S is the current state and A is the proposed action.

This model ensures:

- **Invariant preservation**
- **Non-bypassability**
- **Deterministic evaluation**
- **Execution-time enforcement**

RGA defines the structure of the execution model, the formal semantics establish its correctness conditions, and AGCP enforces these conditions through a deterministic control plane. PBSAI provides the system structure within which these mechanisms operate.

12 Comparative Systems Analysis

Existing approaches to AI governance can be broadly categorized into three classes: policy frameworks, enforcement mechanisms, and runtime systems. While each addresses aspects of governance, none provides end-to-end alignment from intent to execution.

12.1 Policy Frameworks

Frameworks such as the NIST AI Risk Management Framework (NIST, 2023) and the EU AI Act (EU, 2024) define governance objectives, risk categories, and compliance requirements. These frameworks are essential for establishing governance intent but are not directly executable.

As a result, they rely on downstream interpretation and implementation, introducing ambiguity and inconsistency.

12.2 Policy Enforcement Mechanisms

Systems such as policy engines, identity and access management platforms, and guardrails frameworks provide enforcement capabilities. These systems can evaluate rules and restrict actions, but they typically operate in isolation and are tightly coupled to specific platforms.

Furthermore, they often lack deterministic execution semantics and do not guarantee that enforcement decisions are consistently applied at the moment of execution.

12.3 Runtime Systems

Agent frameworks and orchestration platforms enable autonomous execution but frequently lack integrated governance controls. Governance is often externalized, applied inconsistently, or bypassable through alternative execution paths.

12.4 Limitations of Existing Approaches

Across these categories, three systemic limitations emerge:

- **Lack of executability:** Policy intent is not directly translatable into runtime behavior.
- **Fragmented enforcement:** Governance mechanisms are distributed and inconsistent.
- **Absence of deterministic control:** Execution decisions are not guaranteed to be reproducible or auditable.

12.5 Resolution via Integrated Governance Stack

The proposed architecture addresses these limitations:

- Governance compilation enables executable translation of policy intent.
- RGA (Willis, 2026d) provides a canonical execution model for state transitions.
- AGCP (Willis, 2026a) enforces deterministic, execution-bound authorization.
- PBSAI (Willis, 2026e) provides the operational system structure and evidence model.

Together, these components establish a unified governance model in which intent, enforcement, and execution are aligned.

This layered relationship is illustrated in Figure 1, which makes explicit the separation between intent definition, executable translation, runtime semantics, deterministic enforcement, and operational realization.

13 Future Work

Several directions remain for further development. First, empirical validation of the proposed architecture is required, including implementation of AGCP-compliant control planes and evaluation of performance, scalability, and integration with existing enterprise systems.

Second, the formal execution semantics can be extended to address probabilistic and uncertain inputs, particularly in AI-assisted decision pipelines, where correctness must be defined under bounded uncertainty. Additional work is also needed to formalize multi-agent coordination and cross-domain invariant enforcement.

Finally, future work includes the development of reference implementations and interoperability standards, including standardized governance action schemas and validation interfaces as defined in the AGCP Specification, enabling consistent adoption across heterogeneous environments.

14 Conclusion

AI governance is undergoing a fundamental shift from static policy frameworks and post hoc auditing toward execution-time control of autonomous systems. This shift reflects the increasing complexity, autonomy, and risk profile of modern AI-enabled environments.

This paper has shown that existing approaches remain fragmented across policy definition, enforcement mechanisms, and runtime execution. While each contributes valuable capabilities, none provides a unified model that connects governance intent, execution semantics, and enforcement within a single architecture.

To address this gap, we have presented a vertically integrated governance stack. Governance compilation transforms policy intent into executable artifacts. The Runtime Governance Architecture (RGA) (Willis, 2026d) defines governance as a deterministic state-transition execution model. The Artificial Intelligence Governance Control Plane (AGCP) (Willis, 2026a) enforces execution-bound authorization and lifecycle correctness. The formal execution semantics (Willis, 2026c) establish the correctness conditions of this model, including invariant preservation and non-bypassability. The Practitioner's Blueprint for Secure AI (PBSAI) (Willis, 2026e) operationalizes governance within a domain-structured, evidence-centric multi-agent system.

Together, these components establish governance as a control-plane concern, in which decisions are evaluated, enforced, and recorded at the point of execution.

Ultimately, the question is no longer whether governance should be applied, but how it is embedded within the execution fabric of AI systems. This work argues that governance must be treated as a first-class architectural layer, with deterministic execution semantics and control-plane enforcement as its foundation.

References

Bai, Y. et al. Constitutional AI: Harmlessness from AI Feedback. arXiv:2212.08073. https://arxiv.org/abs/2212.08073

CISA. Artificial Intelligence Security Resources. https://www.cisa.gov/ai

Google DeepMind. Strengthening Our Frontier Safety Framework. https://deepmind.google/discover/blog/strengthening-our-frontier-safety-framework/

ENISA. Artificial Intelligence Cybersecurity Challenges. https://www.enisa.europa.eu/publications/artificial-intelligence-cybersecurity-challenges

European Commission. EU Artificial Intelligence Act. https://artificialintelligenceact.eu

Gartner. Global AI Regulations Fuel Billion-Dollar Market for AI Governance Platforms. https://www.gartner.com/en/newsroom/press-releases/2026-02-17-gartner-global-ai-regulations-fuel-billion-dollar-market-for-ai-governance-platforms

Gaurav, S., Heikkonen, J., and Chaudhary, J. Governance-as-a-Service: A Multi-Agent Framework for AI System Compliance and Policy Enforcement. arXiv:2508.18765, 2025. https://arxiv.org/abs/2508.18765

LessWrong. LLM Agents Fail Safety Rules. https://www.lesswrong.com/posts/wRsQowKKbgyXv2eni/i-tested-llm-agents-on-simple-safety-rules-they-failed-in

Mazzocchetti, A. M. Cryptographic Runtime Governance for Autonomous AI Systems: The Aegis Architecture for Verifiable Policy Enforcement. Zenodo, 2026. DOI: 10.5281/zenodo.19027190. https://zenodo.org/doi/10.5281/zenodo.19027190

McKinsey & Company. Superagency in the Workplace: Empowering People to Unlock AI's Full Potential at Work. https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/superagency-in-the-workplace-empowering-people-to-unlock-ais-full-potential-at-work

Meta. Building Generative AI Responsibly. https://ai.meta.com/static-resource/building-generative-ai-responsibly/

Microsoft. Agent Governance Toolkit. https://github.com/microsoft/agent-governance-toolkit

Microsoft. From Runtime Risk to Real-Time Defense: Securing AI Agents. https://www.microsoft.com/en-us/security/blog/2026/01/23/runtime-risk-realtime-defense-securing-ai-agents/

NIST. Artificial Intelligence Risk Management Framework (AI RMF 1.0). DOI: 10.6028/NIST.AI.100-1. https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf

OpenAI. Guardrails Documentation. https://guardrails.openai.com/docs

Reddit. AI Governance Thread. https://www.reddit.com/r/cybersecurity/comments/1iw943r/ai_governance_tools_and_resources/

TechCrunch. Rogue Agents and Shadow AI. https://techcrunch.com/2026/01/19/rogue-agents-and-shadow-ai-why-vcs-are-betting-big-on-ai-security/

Willis, J. (2026a). *AGCP: A Deterministic Execution-Layer Governance Control Plane for Autonomous and Programmatic Systems*. Sustainable Future Tech, Inc. https://doi.org/10.5281/zenodo.19323672

Willis, J. M. (2026b). *Artificial Intelligence Governance Control Plane (AGCP) Specification v0.9.0*. Sustainable Future Tech Publishing. https://doi.org/10.5281/zenodo.18923302 Live repository: https://github.com/jwillisSFT/agcp-spec

Willis, J. M. (2026c). *Formal Execution Semantics and Safety Invariants for Governance Control Planes in Autonomous Systems*. https://doi.org/10.5281/zenodo.19241732

Willis, J. (2026d). *Runtime Governance Architecture: Consistent Governance Execution for Enterprise Systems and Autonomous Agents*. Sustainable Future Tech Publishing. https://doi.org/10.5281/zenodo.19286845

Willis, J. (2026e). *The PBSAI Governance Ecosystem: A Multi-Agent AI Reference Architecture for Securing Enterprise AI Estates*. arXiv:2602.11301. https://arxiv.org/abs/2602.11301

YouTube. The Pre-Execution Control Plane: Why AI Needs Governance Before Action. https://www.youtube.com/watch?v=JIlIn-WCCaw